


# Is It Enough?



We helped MagicWorks™ design their data model then left them to finish data loading... but it hasn't gone well

They need help tuning their data loading processes to get maximum performance from the SQLDW

# Agenda

Loading for performance

Partitioning

Transformations

Load optimizations

Best practices

Loading for performance

# ELT for maximizing scan performance

Do large batch loads

Follow table design best practices

Don't over partition

Use more memory to avoid premature trimming

Manage updates and deletes carefully

Achieve segment elimination by ordering or partitioning

# Do Large Batch Loads

## Target >100,000 per columnstore in each load

With no partitions, this means  $> 100,000 * 60$  rows (~6 million) rows per CTAS or bulk insert

With 4 partitions, this means  $> 100,000 * 60 * 4$  (~24 million) rows per CTAS or bulk insert

## Ideally 1 million per Columnstore in each load

With no partitions, this means  $> 1,000,000 * 60$  rows (~60 million) rows per CTAS or bulk insert

With 4 partitions, this means  $> 1,000,000 * 60 * 4$  (~240 million) rows per CTAS or bulk insert

# Batching trickle loads

Scenario: 100,000 rows per Columnstore, no partitions  
6,000,000+ rows required for each batch load

	500 Rows/Sec	1000 Rows/Sec	2000 Rows/Sec
Load threshold exceeded (hours)	<3.5 hours	< 2 hours	< 1 hour

# Partitioning Guidance

Row groups cannot cross partition boundaries

Over-partitioning impacts row group quality

Over-partitioning impacts compression

Make sure your targeted data set allows for at least 6 million rows per partition. Ideally 60 million rows per partition!

Loading across partitions can impact performance

# Avoid getting trimmed row groups

## Compute memory required for quality row groups

[Memory guidance for Columnstore](#)

Use views provided.

## Grant sufficient memory

Use correct resource class (avoid smallrc)

Scale if needed

Keep load query simple – stage to a Heap/CI if needed

Force serial execution if needed

## Avoid dictionary pressure

Isolate problematic string columns into a separate table if needed



# Stage data with row stores for transformation

## Use heaps or clustered indexes when

- Many modifications required to base data

- Data requires transformation to match target table

## Alter Index Reorganize/Rebuild to defragment

- Reorganize is lighter weight and online

- Rebuild is heavy weight and offline

# Segment Elimination

Columnstore is not ordered

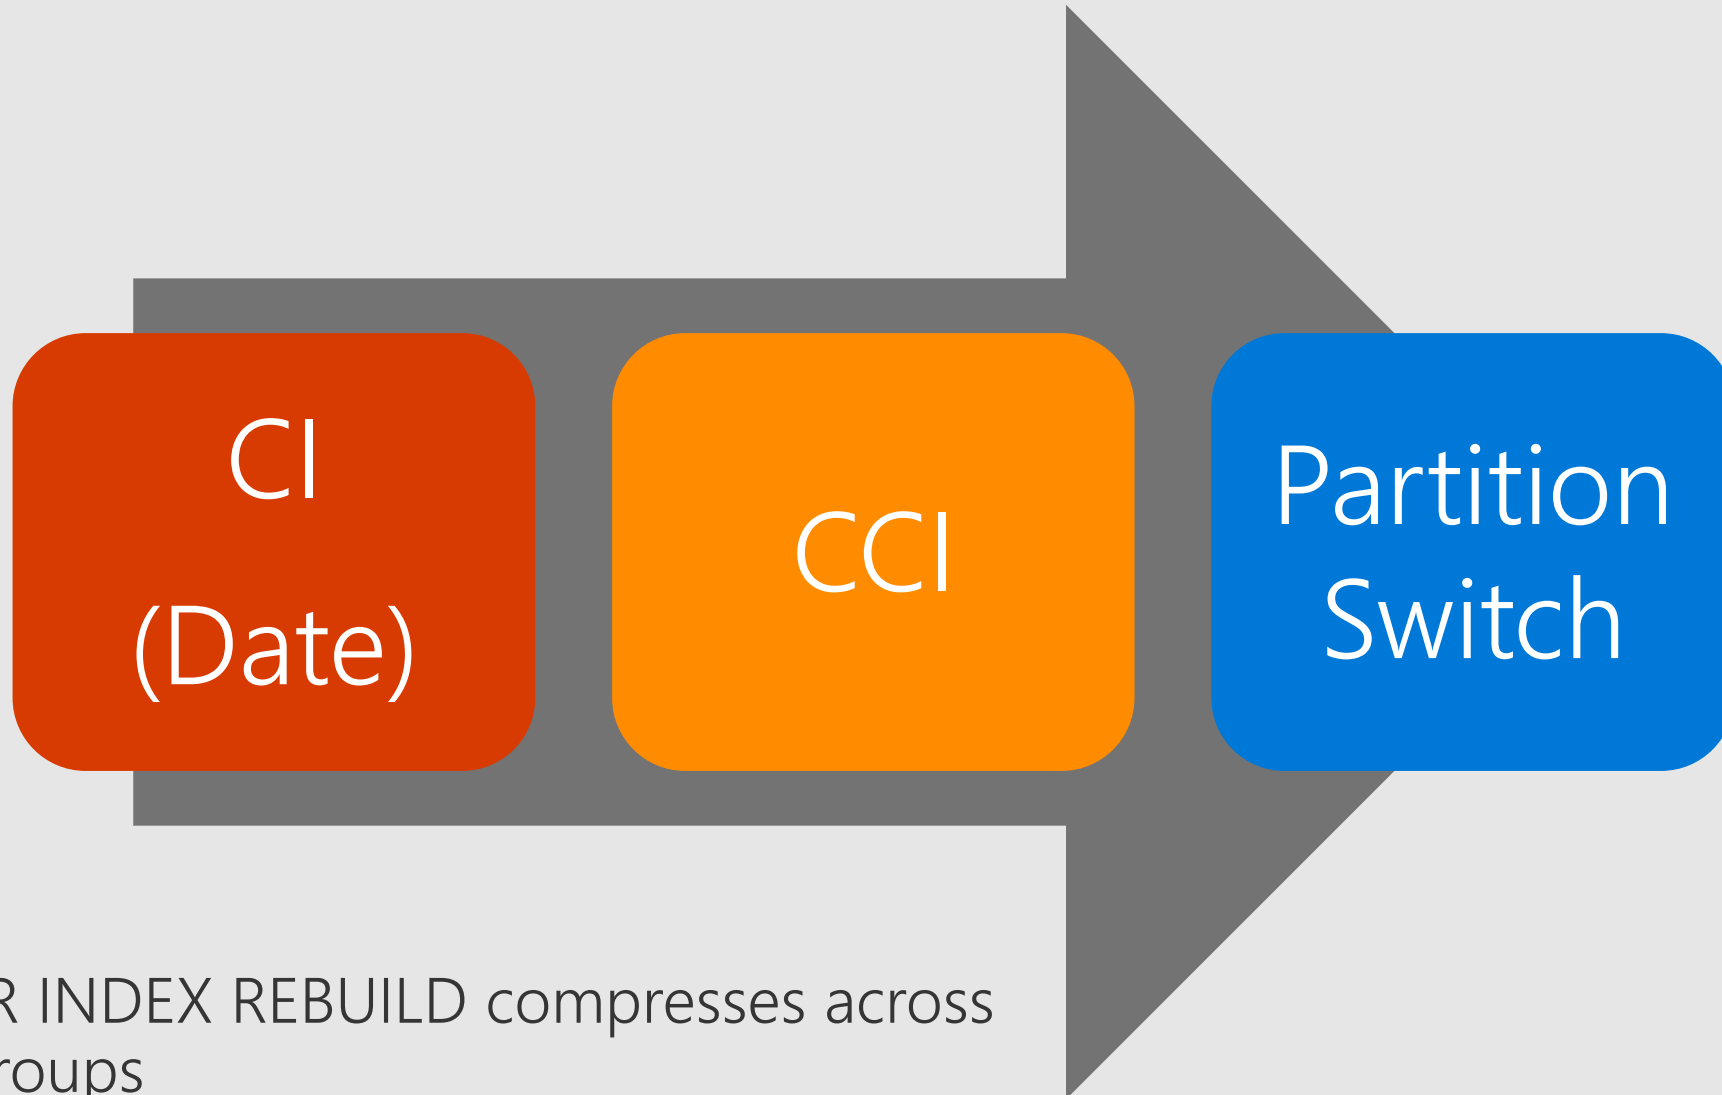
Min,max metadata used to filter out segments

Helps if data arrives naturally ordered e.g. timestamp

Rebuilding indexes will not keep ordering intact

Min	Max
Col 1 val	Col1 val
Col 2 val	Col 2 val
...	...

# Ordered CCI for improved segment elimination



ALTER INDEX REBUILD compresses across rowgroups

Removes any manual ordering from CCI

# Surrogate keys

## Identity $\neq$ Unique

Can have holes and duplicates

## Current Limitations:

Cannot CTAS with Identity:

CREATE TABLE then INSERT...INTO

Identity column cannot be distribution key

# Using the IDENTITY property

```
CREATE TABLE dbo.Dim1
(C1 INT IDENTITY(1,1)
,C2 INT
)
WITH
(DISTRIBUTION = HASH(C2)
,CLUSTERED COLUMNSTORE INDEX
)
;
```

# Data vault key: Use HASHBYTES

```
DECLARE @i VARCHAR(8000) = REPLICATE('X',8000)
```

```
SELECT
```

```
LOWER(CONVERT(CHAR(32),HASHBYTES('MD5',@i),2))
```

```
;
```

```
SELECT
```

```
LOWER(CONVERT(BIGINT,HASHBYTES('MD5',@i),2))
```

```
;
```


# Demo: Managing Surrogates

# Lab 004 - Managing Surrogates



# Table partitioning

# Row & Column Store & Partition



# Why Partition?

## Benefit to Loads

### Data Lifecycle Management

- Drop partition avoids transaction logging

- Insert to empty table/partition avoids transaction logging

- ⇒ Partition Switching pattern

### Targeted Index Builds

## Benefit to Queries

### Partition Elimination

# CREATE TABLE with partitions

```
CREATE TABLE [cso].[FactOnlineSales_PTN]
(
    [OnlineSalesKey]      int          NOT NULL
,   [DateKey]            datetime     NOT NULL
,   [StoreKey]           int          NOT NULL
,   [ProductKey]         int          NOT NULL
,   [CurrencyKey]        int          NOT NULL
,   [SalesQuantity]      int          NOT NULL
,   [SalesAmount]        money        NOT NULL
,   [UnitPrice]          money        NULL
)
WITH
(
    CLUSTERED COLUMNSTORE INDEX
,   DISTRIBUTION = HASH([ProductKey])
,   PARTITION
    (
        [DateKey] RANGE RIGHT FOR VALUES
        (
            '2007-01-01 00:00:00.000', '2008-01-01 00:00:00.000'
        , '2009-01-01 00:00:00.000', '2010-01-01 00:00:00.000'
        )
    )
)
;
```

# Creating a partitioned table with CTAS

```
CREATE TABLE [cso].FactOnlineSales_PTN
WITH
(
    CLUSTERED COLUMNSTORE INDEX
    , DISTRIBUTION = HASH([ProductKey])
    , PARTITION
        (
            [DateKey] RANGE RIGHT FOR VALUES
            (
                '2007-01-01 00:00:00.000' , '2008-01-01 00:00:00.000'
                , '2009-01-01 00:00:00.000' , '2010-01-01 00:00:00.000'
            )
        )
)
AS
SELECT *
FROM [cso].[FactOnlineSales]
;
```

# Partitioning guidance

## Partition for data management

Don't need to use partition elimination for faster performance

## Don't over partition!

Data already spread across 60 distributions

Partitioning granularity likely to differ to SQL Server

Columnstore index row groups give ideal performance when full

Full rowgroups need at least 60 million rows per partition



# Example

Table has 60 billion rows spanning 3 years

Partition granularity	# Rows per distribution	# Partitions	# Rows per partition	# Row groups per partition
Year	1,000,000,000	3	333,333,333	334
Quarter	1,000,000,000	12	83,333,333	84
Month	1,000,000,000	36	27,777,777	28
Day	1,000,000,000	1095	913,242	1
No partitioning	1,000,000,000	1	1,000,000,000	1000

# Transformations





# CTAS

```
CREATE TABLE #Nums
WITH (DISTRIBUTION=REPLICATE, LOCATION=USER_DB)
AS
WITH      L0      AS (SELECT 1 AS c UNION ALL SELECT 1)
,         L1      AS (SELECT 1 AS c FROM L0 AS A, L0 AS B)
,         L2      AS (SELECT 1 AS c FROM L1 AS A, L1 AS B)
,         L3      AS (SELECT 1 AS c FROM L2 AS A, L2 AS B)
,         L4      AS (SELECT 1 AS c FROM L3 AS A, L3 AS B)
,         L5      AS (SELECT 1 AS c FROM L4 AS A, L4 AS B)
,         Nums    AS (SELECT ROW_NUMBER() OVER(ORDER BY c) AS n FROM L5)
SELECT    CAST(n AS BIGINT) as Number
FROM      Nums
WHERE     n BETWEEN @num_Start AND @num_End
OPTION    (LABEL='fn_nums : #nums create')
;
```

Based on fn\_nums by Itzik Ben-Gan

# CTAS vs. Create Table

## CTAS

Table populated

Rowcount set to actual

Data Type from Select

Nullability from Select

Default constraints: no

## CREATE TABLE

Table is empty

Rowcount set to 1000

Data types from DDL

Nullability from DDL

Default constraints: yes

# CTAS vs. INSERT...SELECT

## CTAS

New object created

No column statistics on table

Row count updated

Page count updated

## INSERT..SELECT

Existing object populated

Any stats maintained – UPDATE

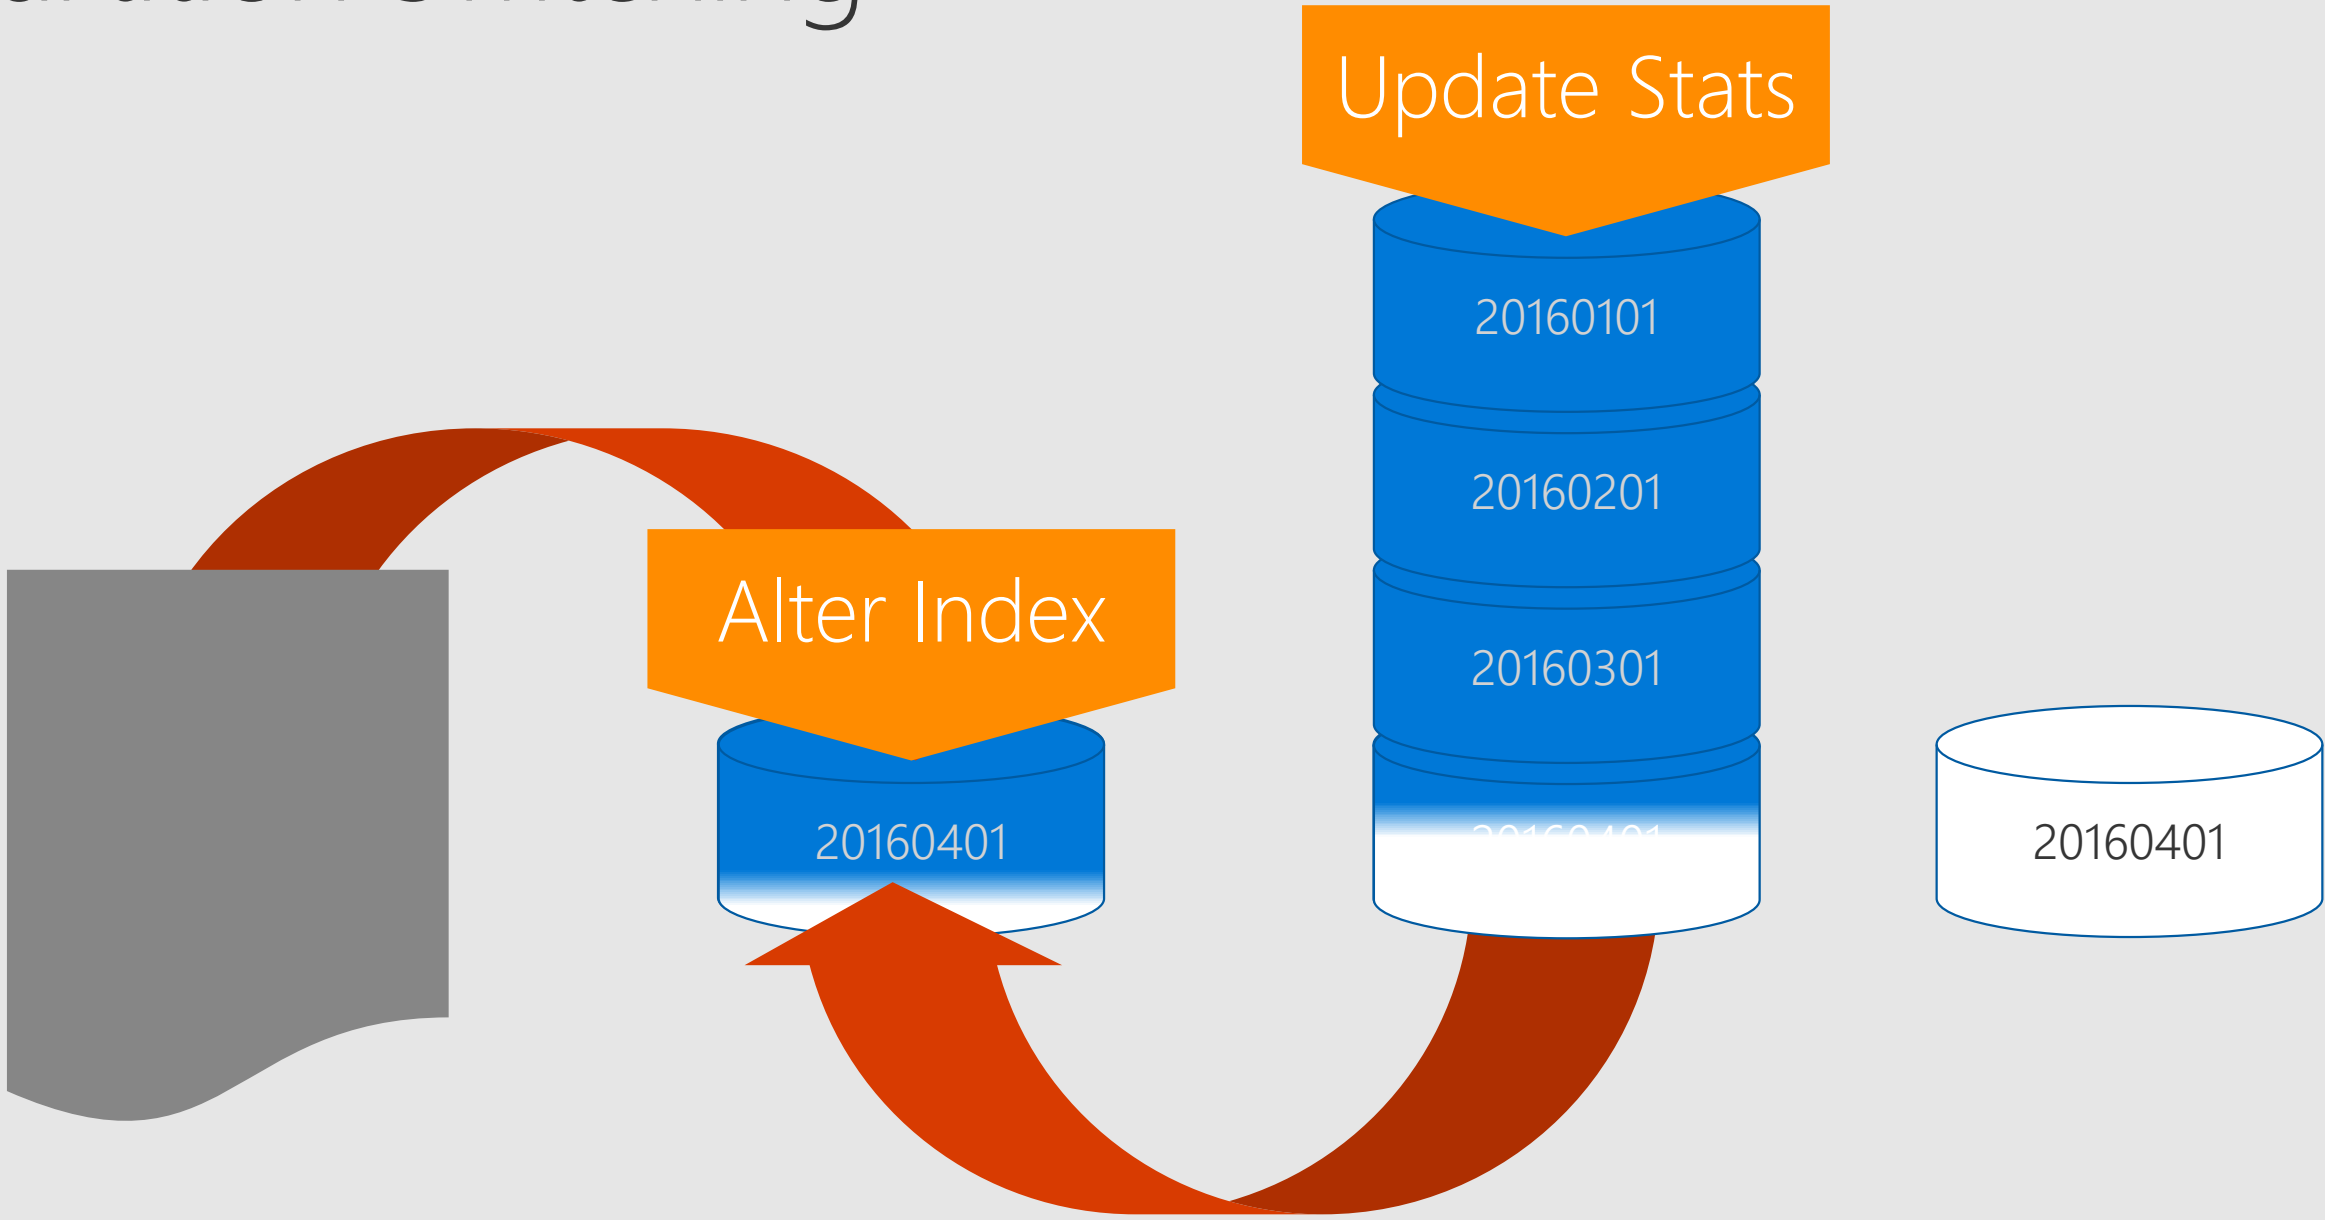
Row count not updated

Page count not updated

# UPSERTing data with CTAS

```
CREATE TABLE [tmp].[DimProduct]
WITH (DISTRIBUTION = ROUND_ROBIN)
AS -- New rows and new versions of rows
SELECT      s.[ProductKey]
,           s.[ProductName]
,           s.[ColorName]
FROM        [src].[DimProduct] s
UNION ALL --Keep rows that are not being updated
SELECT      p.[ProductKey]
,           p.[ProductName]
,           p.[ColorName]
FROM        [cso].[DimProduct] p
WHERE NOT EXISTS
(   SELECT *
    FROM [src].[DimProduct] s
    WHERE s.[ProductKey] = p.[ProductKey]
)
;
```

# Partition Switching



# CTAS – Create partition for switch out

```
CREATE TABLE [cso].[FactOnlineSales_out]
WITH
(
    DISTRIBUTION=HASH ([ProductKey])
,   CLUSTERED COLUMNSTORE INDEX
,   PARTITION ([DateKey]
    RANGE RIGHT FOR VALUES ( '2007-01-01 00:00:00.000'
                              , '2008-01-01 00:00:00.000'
                              )
    )
)
AS
SELECT *
FROM    [cso].[FactOnlineSales]
WHERE 1=2;
```

# CTAS – Create Partition for Switch In

```
CREATE TABLE [cso].[FactOnlineSales_in]
WITH
(
    DISTRIBUTION=HASH ([ProductKey])
,   CLUSTERED COLUMNSTORE INDEX
,   PARTITION ([DateKey]
    RANGE RIGHT FOR VALUES ('2007-01-01 00:00:00.000', '2008-01-01 00:00:00.000'
    )
    )
)
AS
SELECT *
FROM [cso].[FactOnlineSales_ptn] tgt
WHERE tgt.[DateKey] >= '2007-01-01 00:00:00.000'
AND    tgt.[DateKey] <  '2008-01-01 00:00:00.000'
UNION ALL
SELECT *
FROM [cso].[FactOnlineSales] stg
WHERE stg.[DateKey] >= '2007-01-01 00:00:00.000'
AND    stg.[DateKey] <  '2008-01-01 00:00:00.000'
;
```



# Perform the switches

```
ALTER TABLE [cso].[FactOnlineSales_ptn]  
SWITCH PARTITION 2  
TO [cso].[FactOnlineSales_out] PARTITION 2  
;  
ALTER TABLE [cso].[FactOnlineSales_in]  
SWITCH PARTITION 2  
TO [cso].[FactOnlineSales_ptn] PARTITION 2  
;
```



0 - 20
20 - 40
40 - 60
60 - 80
80 - 100
100 - 120
> 120



40 - 60



1. UNION New/Updated  
and Unchanged Records

CTAS into New Table



## 2. Switch Out Old Partition


```
ALTER TABLE dbo.Fact SWITCH PARTITION 3 TO  
dbo.Fact_out;
```



### 3. Switch In New Partition

```
ALTER TABLE dbo.Fact_new SWITCH TO  dbo.Fact  
PARTITION 3;
```

# Demo: Partitioning Scripts



Each data movement is performed by a stored procedure that follows a standard pattern, using a set of utility procedures

Begin Audit

CTAS

Gather Rowcount

Generate Stats

Finalise Audit

# Load optimizations

# Factors improving performance

Index choice

System scale

Distribution type


File format type

File system layout

Row number in file




# Index type matters




# Scaling matters

DWU	Max External Readers	Max Writers CCI / Heap	Max Writers CI / NCI
100	8	60	60
200	16	60	60
300	24	60	60
400	32	60	60
500	40	60	60
600	48	60	60
1000	80	60	60
1200	96	60	60
1500	120	120	60
2000	160	120	60
3000	240	240	60
6000	480	480	60


# Scale matters



# Table distribution matters



# File format matters (delimited text)



# Delimited text guidance

Evenly split the data into multiple files


One file per reader

Delimited text is the fastest

Compressed text limits  
concurrent access to  
text files

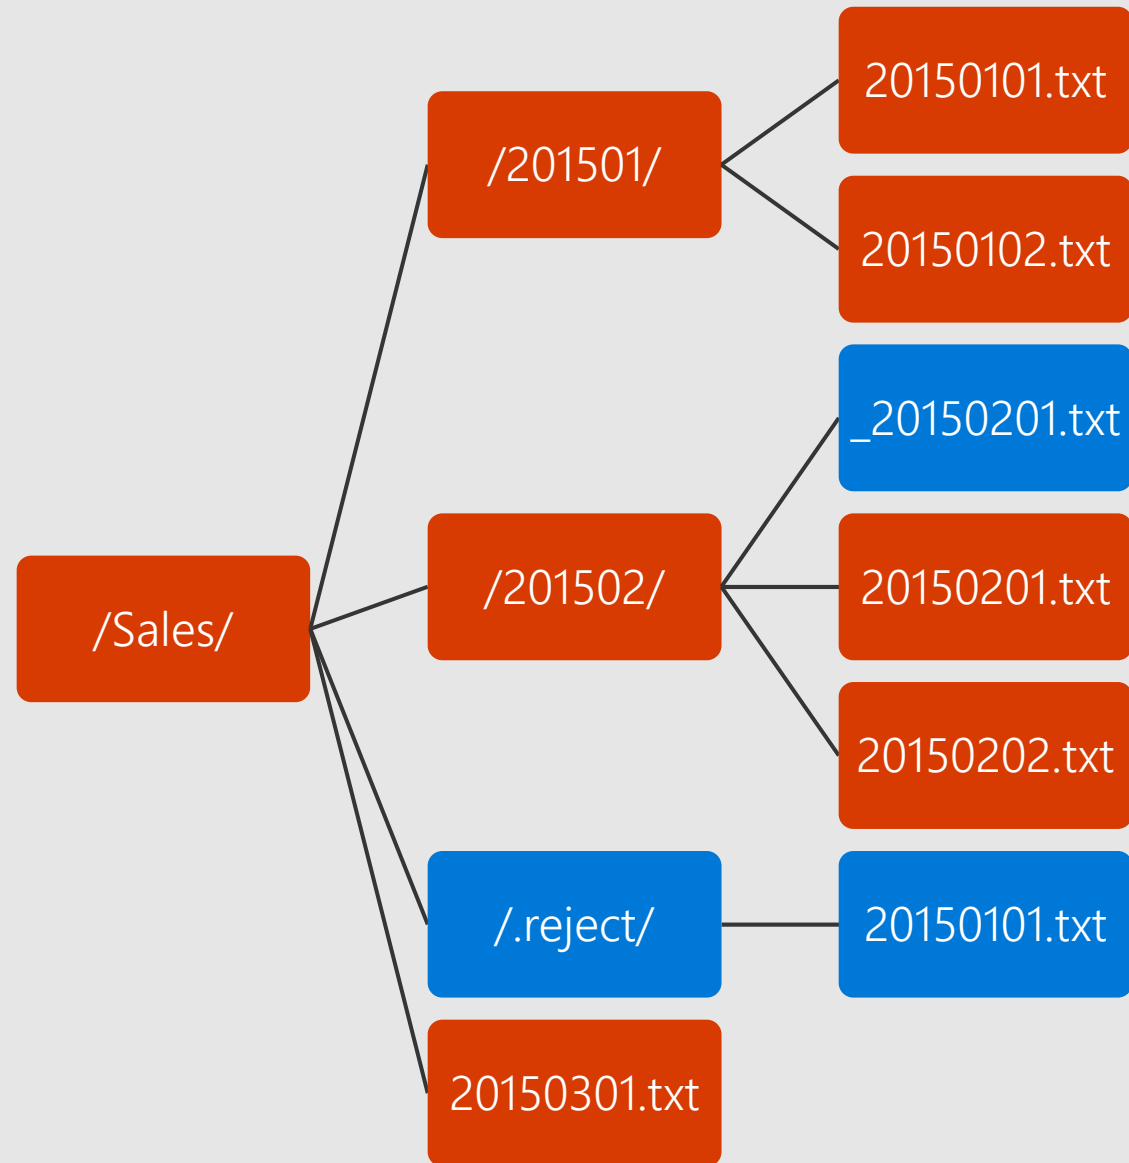
Split data across files  
OR  
use different file format

# File format matters (Big Data)



# File system layout

Objects prefixed  
with \_ and . are  
ignored recursively





# Optimised Source file

```
RowNumber,Operation,ProductKey,EnglishProductName,Color
1,I,10000,GriffyndorAreMagic,Red
2,U,10100,RavenclawAreStillQuiteGood,Yellow
3,D,10200,SlytherinPleaseFinishThird,White
4,I,10300,CongratulationsHufflepuff,Blue
```

# Upsert Logic – CTAS optimised

```
CREATE TABLE tmp.DimProduct
WITH (DISTRIBUTION = ROUND_ROBIN)
AS -- New rows and new versions of rows
SELECT      s.ProductKey
,           s.ProductName
,           s.ColorName
FROM        [src].[DimProduct] s
WHERE       s.Operation IN ('I', 'U')
UNION ALL --Keep rows that are not being updated
SELECT      p.ProductKey
,           p.ProductName
,           p.ColorName
FROM        [cso].[DimProduct] p
WHERE NOT EXISTS
(   SELECT *
    FROM [src].[DimProduct] s
    WHERE s.ProductKey = p.ProductKey
);
```

# Best Practices

# Loading takeaways

## Index Choice

Data shows not much difference between CCI and heap tables

Network is the bottleneck

## DWU

Increasing scale of system automatically increases load performance

Particularly evident at the low-end

## Distribution type

Round Robin fastest

Poor distributions can result in bad performance

# Loading take-aways

## File type

UTF-16 faster due to less time converting encoding

## Resource class

Matters most for column store segment health than load speed

# Dimension tables

## Use round robin for small tables

Perform all DML before inserting into Replicated table

## Use clustered index for ordered scans

Clustered columnstore index is not ordered

Small tables (<60M) may benefit more from being row stores

## Load in full where possible

Dimensions are typically small and are typically mastered in MDM system

## Use metadata rename to reload data

Keeps old and new versions available for easy comparison

# Fact tables

## Land data from ASB in a staging table

- Adheres to an ELT pattern

- Greater control in T-SQL

## Partitions reduce loading impact on production tables

- Partition switch is a metadata operation

## ASB directory structure can scope load

- Improves performance

- Limits compute utilization

# Applying Type 4 Partitioning





# Summary

# Summary

Do not load data via singleton inserts

Very inefficient

Use PolyBase for batch loading

Land data in blob storage first

Use stored procedures to maximize control

Adhere to an ELT load pattern

Trickle load/micro batch

Use bcp if PolyBase isn't an option

# So What?

# MagicWorks

Implement ELT Loading Scripts to:

- Use CTAS

- Partition Switch

- Index Design

Data Loading Script

Structure flat file storage effectively